

AInimotion

Layer-Separated Deformable Interpolation for Anime Video

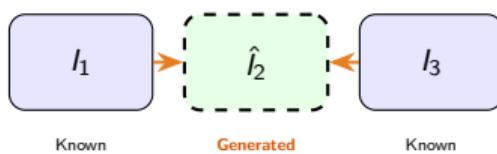
Travis Whitney

February 19, 2026

Video Frame Interpolation for Anime

Why anime is uniquely challenging:

- ① **Flat color regions** — no texture for optical flow
- ② **Non-linear motion** — characters “pop” between poses
- ③ **Ink line preservation** — thin lines blur easily
- ④ **Scene cuts** — abrupt shot changes
- ⑤ **Layer separation** — backgrounds pan rigidly, characters move independently



Why Deep Learning?

Frames have **multi-scale spatial structure** (edges, textures, objects) and **temporal structure** (motion patterns across frames). Anime adds **layer structure**: rigid backgrounds vs. deformable characters. My model encodes these structural priors directly.

Prior Work & My Contribution

Prior Approaches:

Method	Idea	Gap
DAIN [7]	Depth-aware flow	Not anime-specific
AdaCoF [2]	Deform. kernels	Single motion path
AnimeInterp [1]	Segment match	Ext. segmentation
RIFE [3]	Lightweight flow	Blurs heavy motion
LDMVFI [8]	Latent diffusion	100× slower

Layer decomposition dates to Wang & Adelson (1994) [6], but **no prior VFI method** builds anime's layer structure into the architecture.

What makes Alnimotion novel:

- ➊ **Dual-path architecture** — dedicated BG (affine grid) + FG (AdaCoF kernels)
- ➋ **Domain-informed** — encodes how anime is produced (cel layers)
- ➌ **Learned compositor** — soft α mask, no external segmentation
- ➍ **Edge-aware loss** — Sobel-weighted L1 protects ink lines
- ➎ **Practical** — real-time inference (~50ms) vs. diffusion (~5–30s)

ATD-12K Dataset

Source: ATD-12K [1] (CVPR 2021)

- **12,000** frame triplets (I_1, I_2, I_3) from **30 films**
- Motion categories: small, medium, **large**
- Established anime VFI benchmark

Training setup:

- Random 384×384 crops
- Horizontal/vertical/temporal flips
- 100K samples per epoch

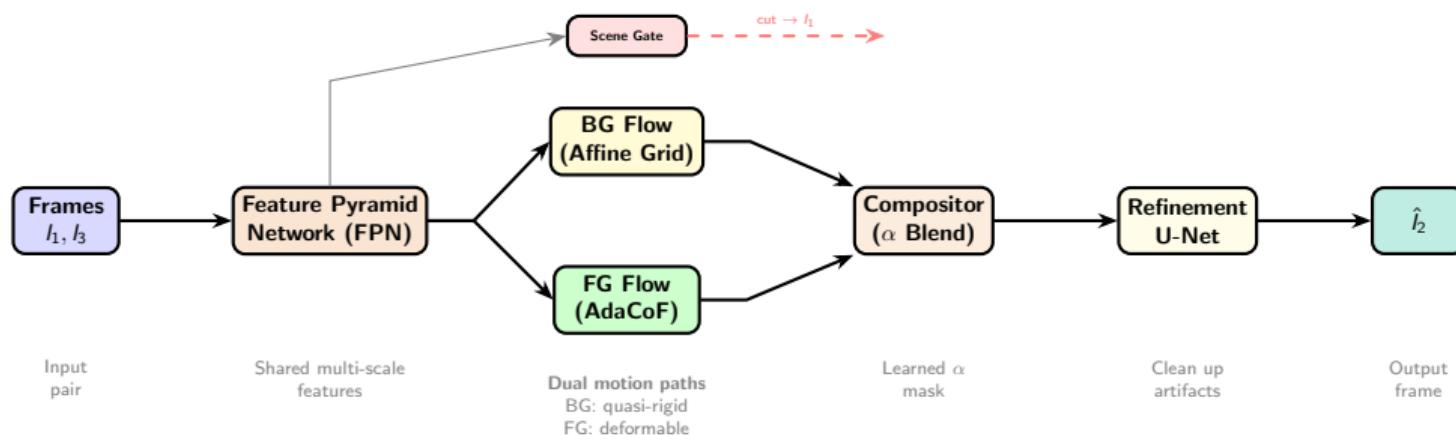
Why This Dataset?

- Purpose-built for anime VFI
- Publicly available
- Rich motion diversity
- Includes difficulty labels

Triplet Format

(I_1, I_2, I_3) : I_2 is ground truth used as supervision during training.

Architecture Overview: LayeredInterpolator



Key Components & Loss Functions

Components:

- **FPN:** Shared encoder, 4 scales ($\frac{1}{2}$ to $\frac{1}{16}$) + **correlation volumes** (per-pixel dot-product similarity over a 9×9 search window between frames — tells the network “where did each pixel move?”)
- **Scene Gate:** Detects hard cuts via correlation statistics, bypasses interpolation
- **Background:** 8×8 quasi-rigid affine grid for camera pans/zooms
- **Foreground (AdaCoF):** 9×9 deformable sampling kernels per pixel (81 weights + 162 offsets = 243 params/pixel)
- **Composer:** Learned soft α mask + refinement U-Net

Generator Loss:

$$\mathcal{L}_G = \underbrace{1.5 \cdot \mathcal{L}_{L1}}_{\text{Reconstruction}} + \underbrace{0.1 \cdot \mathcal{L}_{perc}}_{\text{VGG19}} + \underbrace{1.0 \cdot \mathcal{L}_{edge}}_{\text{Ink Lines}} + \underbrace{0.005 \cdot \mathcal{L}_{GAN}}_{\text{Phase 2}}$$

- \mathcal{L}_{L1} : Pixel accuracy, drives PSNR
- \mathcal{L}_{perc} : VGG19 feature similarity, preserves style
- \mathcal{L}_{edge} : Sobel-weighted L1, 20 \times multiplier on edges
- \mathcal{L}_{GAN} : LSGAN patch discriminator with label smoothing (0.1), Phase 2 only

Two-Phase Training Strategy

Phase 1: Reconstruction (Ep 0–34)

- Generator only — no adversarial loss
- $\text{lr}_G = 3 \times 10^{-4}$
- Loss: $\mathcal{L}_{L1} + \mathcal{L}_{\text{perc}} + \mathcal{L}_{\text{edge}}$
- **Goal:** PSNR ≥ 28 dB

$\Downarrow D$ warmup (500 batches)

Phase 2: GAN Fine-Tuning (Ep 35–49)

- Adversarial loss activated ($\lambda_{\text{GAN}} = 0.005$)
- $\text{lr}_G = 5 \times 10^{-5}$ (reduced)
- D trains every other batch (1:2 ratio)
- **Goal:** perceptual sharpness

Why two phases?

- ① Generator learns stable reconstructions *before* seeing adversarial gradients
- ② GAN adds sharpness without destabilizing a converged generator

Key stabilization choices:

- **D warmup** — 500 batches before G sees GAN loss
- **Low GAN weight** (0.005) prevents mode collapse
- **Adaptive lr_D** — auto-adjusts based on d_{acc}
- **Label smoothing** (0.1) on discriminator

GAN Stabilization & Training Infrastructure

Discriminator Control:

- **D Warmup** (500 batches) — D learns before G sees adversarial loss
- **Patch Discriminator** — 70×70 patches, not full images
- **Label Smoothing** (0.1) — prevents overconfident D
- **D Update Ratio** (1:2) — D trains every other batch
- **Adaptive lr_D** — auto-adjusts based on d_{acc}

Evaluation Metrics:

- **PSNR** (dB): $-10 \log_{10}(\text{MSE})$. Higher = better pixel accuracy. Log scale: +3 dB \approx halving MSE.
- **SSIM**: Measures luminance, contrast, and structural similarity (0–1). Catches distortions PSNR misses.
- **Visual comparison**: Side-by-side with ground truth, especially on ink lines and motion.

Infrastructure

Mixed precision, grad_clip=1.0, OOM recovery, W&B logging. RTX 5090 (32 GB VRAM).

Design Evolution & Training Progress

Component	Initial	Final
Base Channels	32	96 (3×)
Kernel Size	7	9 (sweep winner)
GAN Training	From epoch 0	Two-phase (ep 35)
L1 Weight	1.0	1.5 (sweep-optimized)
GAN Weight	0.01	0.005 (gentler)
Edge Multiplier	10×	20× (sharper)
Disc. LR	10^{-4}	2×10^{-5}
D Update	1:1	1:2 (G trains more)
D Warmup	None	500 batches

Lesson: GAN training is highly sensitive in anime VFI. The discriminator easily dominates.

Evaluation Baselines (ATD-12K):

Method	PSNR	SSIM
RIFE	~25–27 dB	~0.85
Animelinterp	~27–29 dB	~0.88
Alnimotion (target)	≥ 28 dB	≥ 0.90

Sweep Results (13 experiments):

- Best: **21.80 dB** ($K=9$, crop 384, perc=0.1)
- Crop 384 alone: +0.74 dB over 256
- Gradients stable ($\sim 0.2\text{--}0.5$)

Targets

Phase 1: PSNR ≥ 28 dB by epoch 34

Phase 2: + perceptual sharpness via GAN

Training Iterations & Experimentation

Phase	Config	Key Insight
Prototype	32ch, K=7	Gradient explosion → added clipping
Scale-up	48ch, K=7	28.7 dB test; D dominance issue
Probes (4)	48ch varied	D weakening helps but not enough
Extensions (5)	32ch/48ch	Higher GAN weight → instability
Phase 3	64ch, K=7	Edge-aware + adaptive lr_d
GPU Upgrade	5070 Ti → 5090	2× VRAM (16→32 GB)
Sweep (13)	96ch varied	K=9, crop 384, perc 0.1 wins
Final	96ch, K=9	In progress (ep 7+)

What changed and why:

- **4 capacity levels** (32→48→64→96 channels)
- **9 GAN balance experiments** to tame discriminator
- **13-experiment sweep** covering lr, loss weights, kernel, crop size
- **Gradient accumulation** for large crops on 32 GB VRAM

Total Experiments

30+ configurations tested
across 7 major training phases

Summary & Next Steps

What I built:

- **LayeredInterpolator** — novel dual-path architecture separating BG/FG motion
- **Edge-aware losses** for ink line preservation
- **Two-phase training** with D warmup for stable GAN integration
- **Robust infrastructure** — OOM recovery, adaptive GAN balancing, mixed precision, W&B logging

Next steps:

- Complete full 50-epoch training run
- Evaluate Phase 2 GAN fine-tuning
- Qualitative evaluation on test set
- Compare with RIFE, AnimelInterp baselines

Project Repository

github.com/TWhit229/AInimotion

References I

-  L. Siyao, S. Zhao, W. Yu, W. Sun, D. Metaxas, C. Loy, Z. Liu.
“Deep Animation Video Interpolation in the Wild.”
CVPR, 2021.
-  H. Lee, T. Kim, T. Chung, D. Pak, Y. Ban, S. Lee.
“AdaCoF: Adaptive Collaboration of Flows for Video Frame Interpolation.”
CVPR, 2020.
-  Z. Huang, T. Zhang, W. Heng, B. Shi, S. Zhou.
“Real-Time Intermediate Flow Estimation for Video Frame Interpolation.”
ECCV, 2022.
-  I. Goodfellow et al.
“Generative Adversarial Nets.”
NeurIPS, 2014.

References II

-  X. Mao, Q. Li, H. Xie, R. Lau, Z. Wang, S. Smolley.
“Least Squares Generative Adversarial Networks.”
ICCV, 2017.
-  J. Wang, E. Adelson.
“Representing Moving Images with Layers.”
IEEE Trans. Image Processing, 1994.
-  W. Bao, W. Lai, C. Ma, X. Zhang, Z. Gao, M. Yang.
“Depth-Aware Video Frame Interpolation.”
CVPR, 2019.
-  D. Danier, F. Zhang, D. Bull.
“LDMVFI: Video Frame Interpolation with Latent Diffusion Models.”
AAAI, 2023.